

# Data Science in football: Comparing Model Predictions with Betting Market Odds

## # Abstract

This project aims to predict football match outcomes by estimating the probabilities of each outcome possible (Home win, Draw, Away win) and comparing them to the implied probabilities of the bookmakers. These implied probabilities are obtained by averaging the odds from 3 selected bookmakers (Pinnacle, Bet365 and William Hill) coming from Football-data.co.uk and then removing the bookmakers' margin. After collecting, cleaning, and merging the datasets, I construct several engineered features (e.g. rolling\_xg\_for\_5, elo\_home, etc.) using football metrics from FBref.com. Once all the datasets are ready, I train 3 different machine learning models (Linear Regression, Random Forest and XGBoost) using an 80/20 train-test split. After training, the models are calibrated using Platt scaling and Isotonic regression. Based on the lowest log-loss value (0,907707), the best-performing model is the raw XGBoost model. RMSE and MAE are used to evaluate the model performance for each outcome possible. To better understand the discrepancies between my model's probabilities and those of the average bookmaker, I analyze systematic biases such as home advantage and team-level deviations between model-implied and bookmaker probabilities. Finally, calibration plots, heatmaps, and scatterplots are produced to provide clear visual insights, ensuring that the results remain easy to interpret.

**Keywords:** data science, Python, machine learning, football-data, prediction-model, data-visualization

## **Table of Contents**

1. Introduction
2. Literature Review
3. Methodology
4. Results
5. Discussion
6. Conclusion
7. References
8. Appendices

## 1. Introduction

The football of the 21st century has become a place where data is king. Almost everything that happens on a football pitch can be tracked. Most clubs now have entire team dedicated to data analysis, both for player-related data and team-level performance metrics. But they are not the only ones that massively use these data. In fact, bookmakers also rely heavily on them, in addition to real-world information to give probabilities to each of the outcomes possible (home win, draw, away win) as well as many special bets (over/under 3.5 goals etc.). Once the bookmakers add their margin, the odds are created. As a huge football fan and former bettor, I'm always fascinated by how the odds for big team can be so high and so "small" for smaller teams and vice versa.

The focus in this project is the following: to what extent can football performance metrics (such as expected goal (xG), recent strength etc.) be used to predict match outcomes probabilities using machine learning models and how these model-implied probabilities differ from bookmakers' odds? The analysis relies on datasets collected from Football-Data.co.uk and advanced football metrics from FBref.com.

The first goal of the project is to build a strong predictive model for football outcomes by comparing different machine learning algorithms such as Linear regression, Random Forest and XGBoost. The second goal is to analyze how, where and why differences arise between the model's predictions and the implied probabilities of the average bookmaker.

The structure of the report is as follows. Section 2 presents the relevant literature. Section 3 describes the methodology: how the datasets are cleaned and merged, the feature engineering procedures, the models used, as well as the calibration and model evaluation process. Section 4 reports the results of the different models. Section 5 discusses the implications of these findings and potential biases such as home advantage and others. Section 6 summarize the findings and provides suggestions for future improvements.

## 2. Literature Review

Back in the day, the first football predictions models relied on the Poisson-based assumptions for goal scoring, with the aim of deducting match outcomes from scoreline probabilities. In last decades, the expected goal (xG) has emerged as a key metrics indicator to evaluate the performance of a team. xG is better than previous statistics like number of shots or percentage of possession to show real performance of a team (Rathke, 2017). They improve also the results of outcomes predictions models according to Mead , Pollard and Reep (2023). Alongside advanced metrics, Hvattum & Arntzen (2010) have shown that opponent-adjusted measures such as Elo ratings (coming from the chess world) is also a precious tool for predicting the real level of a team as the model is also based on the strength/weakness of the opponent. Another useful thing is the use of mov-

ing average (rolling performance feature), indeed this help to teach the model the recent form of a team. As many research have proved, machine learning algorithms, especially tree-based models, have shown to be more accurate than linear models at result predictions (Tax and Joustra, 2015). Nevertheless, these algorithms predicted probabilities suffer from poor calibration. Niculescu-Mizil and Caruana (2005), tells us that calibration techniques, like Platt scaling and Isotonic regression, must be employed to have proper probabilistic prediction. When comparing predicted model outcomes with bookmakers' odds, it should be kept in mind that market structural bias needs to be accounted for because bookmakers make use of it. One of the strongest biases according to Levitt (2004) is that bettors tend to over bet the underdogs which result in weaker odds for underdogs in comparison to their true probability. As the odds set by any bookmaker include a margin, it's necessary to extract the unbiased implied probabilities (Clarke, 2017). This project builds on these findings by comparing calibrated model probabilities with margin-adjusted bookmaker probabilities, with added analysis of how, where, and why divergences arise.

### 3. Methodology

#### 3.1 Data Description

This project is based on two main sources of data concerning the 2022-2023 English premier league season. The first source is Football-data.co.uk. It allows us to collect the odds of each game for more than ten different bookmakers and the important information to identify a game (date, time, names of both teams). There are 380 lines and 106 columns in this dataset. The second was obtained from a publicly available dataset extracted from FBref.com and published on GitHub by the user Dataquestio. This dataset allows us to collect the important metrics for the project (xG for, xG against etc.). However further work was required to standardize team names, remove inconsistencies, and ensure mergeability with bookmaker data. There are 4788 lines and 26 columns in this dataset. After the cleaning, merging and engineering of features, the final dataset (features\_match\_long\_elos\_22\_23.csv) used for the analysis consist of 760 lines (38 games in a season \* 20 teams) and 33 columns. The dataset is fully tabular, combining both numerical and categorical variables (home team, away team, etc.) with each row representing a football match from a team perspective (2 rows for one game) and all variables describing pre-match information. Nine features were created. The most important one being elo\_diff\_for\_team. The Elo rating is a pre-match-strength indicators based on past results. Rolling metrics such as "rolling\_xg\_for\_5", "rolling\_xg\_against\_5" were also a big part. They were no missing value that could affect the proper functioning of the analysis.

### 3.2 Approach

Three models were used to predict the implied probability of the outcomes and therefore the results: 1. Linear regression: This is a classification algorithm that is good for predicting probability. It estimates the probability of a binary or multi-class outcome, as in our case, Win/Draw/Loss. Quick to train and easy to understand, it's often a very good baseline model. 2. Random Forest: It's a strong nonlinear model that handles well diverse football features (xG, rolling metrics, Elo). It's very efficient to capture nonlinear and complex interactions. For example, recent form may have an impact only if the team plays away. However, we must be careful with the overfitting issues 3. XGBoost: This model builds trees sequentially, each new tree tries to correct the residual mistakes of all the preceding trees. It's one of the most powerful techniques for structured data. In general, it's the most accurate model when dealing with probabilistic predictions. However, XGBoost requires many hyperparameters that are not always easy to tune or to understand. Before running the algorithms, a preparation phase was needed. First, the two datasets were cleaned by selecting only the necessary columns for the 2022-2023 season and by normalizing team names and date format across all datasets because several naming variations existed for a team, sometime even within the same dataset (e.g. from "Man City" to "Manchester City"). New variables were created such as "prob\_book\_home", which is the average of the implied probability from bookmakers once the margin is removed. Once the columns names were standardized everywhere, an inner merge could be done using the keys: "date", "home\_team" and "away\_team". After that, feature engineering could begin. Rolling metrics were the biggest part of the work, they consist of averaging the last 5-games values of one statistic to capture the recent form of a team. Finally, an Elo rating was created and Elo-based features were added to the final dataset. All 3 models follow the same architecture: First, load and prepare the datasets (NaN had to be removed, NaN were present in the rolling features) then sort and split the datasets using a chronological train-test ratio. This ratio of 0.8 permits the model to train on the first 80% of the dataset and after train on the last 20% (the sorting phase is important to avoid information leakage). After encoding the target (H/D/A) and scaling the features (only needed for the baseline model) the training of the model can be achieved. It's important to try different values of hyperparameter to reduce overfitting. Once the model is trained, the results are evaluated using 3 tools: accurate, log-loss and Mean brier score. A calibration curve is generated and saved in the result folder to assess how close the predicted probabilities are to the frequencies. A calibration experiment is also realized to compare raw probabilities with Platt scaling and Isotonic regression. The comparison to the bookmakers is made once the 3 models are running correctly and the best one is selected (based on the smallest log-loss). The comparison is done using RMSE and MAE and several scatter plot, heatmap and barplot are drawn trying to explain why and how the model's probabilities differ from the bookmakers.

### 3.3 Implementation

The project is fully implemented in Python, using a set of widely used scientific and machine-learning libraries. The main packages are: Numpy, Pandas, Matplotlib, Seaborn, scikit-learn, XGBoost, pathlib. They are useful for data manipulation, visualization, modeling, evaluation metrics etc. The codebase is organized into a set of modular scripts stored in the /src folder, each responsible for one specific step of the pipeline. This structure brings clarity and reproducibility and makes the project easier to maintain. The main file is main.py, which orchestrates the full workflow by calling the functions defined in each script. The key code components are the first three modules who clean and merge the two dataset, because without that nothing can be achieved. Last two are also important as they provide the answers to the research question.

Example code snippet:

```
# Add rolling xG features over the last N games per team.
def add_rolling_xg_features(df: pd.DataFrame, window: int = 5) -> pd.DataFrame:
    df = df.copy()                      # Work on a copy to avoid modifying the original DataFrame
    df["rolling_xg_for_5"] = np.nan       # Initialize new columns with NaN
    df["rolling_xg_against_5"] = np.nan
    df["rolling_xg_diff_5"] = np.nan

    for team in df["team"].unique():           # Process each team separately
        team_mask = df["team"] == team

        # xG for (offensive form)
        team_xg_for = df.loc[team_mask, "xg_for"]      # Select only the xG for values
                                                       # corresponding to the current team
        df.loc[team_mask, "rolling_xg_for_5"] = (
            team_xg_for.shift(1).rolling(window).mean())
        # .shift() to shift the series by 1 so the current game doesn't use its own
        # xG value (prevents data leakage).
        # .rolling() create a rolling window over the last 5 past matches.
        # .mean() compute the average xG over the rolling window.

        # xG against (defensive form)
        team_xg_against = df.loc[team_mask, "xg_against"]
        df.loc[team_mask, "rolling_xg_against_5"] = (
            team_xg_against.shift(1).rolling(window).mean())

        # xG diff
        team_xg_diff = df.loc[team_mask, "xg_diff"]
        df.loc[team_mask, "rolling_xg_diff_5"] = (
            team_xg_diff.shift(1).rolling(window).mean())
    return df                                # Return the updated DataFrame with rolling features added
```

## 4. Results

### 4.1 Experimental Setup

Hardware: Experiments were executed on a standard consumer laptop (HP), powered with an Intel Core i7 processor. No dedicated GPU was used; the models were trained on a standard CPU. Software: Python 3.12.6, Pandas 2.3.3, seaborn 0.13.2, xgboost 3.1.2, matplotlib 3.10.7, scikit-learn 1.7.2 This setup is lightweight and fully reproducible on any standard machine. Machine learning models include several adjustable settings, called hyperparameters. These settings influence how the model learns and how complex it can become. They act like “instruction rules” that help for different things such as to avoid the overfitting issue and to improve the stability. These values were selected after trying multiple combinations “by hand” and look which one provides the smallest log-loss and smallest overfit. For logistic regression, the main element is the optimization method (“lbfgs”), which ensures that the model will converge correctly even with many features. For Random Forest, the most important hyperparameters control the structure of the trees: n\_estimators=400, max\_depth=6, min\_samples\_split=5, min\_samples\_leaf=4, max\_features=“sqrt”, random\_state=42, n\_jobs=-1, bootstrap=True, max\_samples=0.8. For XGBoost, some additional regularization hyperparameters were necessary because the model is more flexible. objective=“multi:softprob”, num\_class=3, n\_estimators=200, learning\_rate=0.05, max\_depth=3, min\_child\_weight=6, gamma=1, subsample=0.7, colsample\_bytree=0.7, reg\_lambda=3.0, reg\_alpha=1, random\_state=42, n\_jobs=-1, eval\_metric=“mlogloss”,

Since the models used in this project are tree-based (Random Forest, XGBoost) and not neural networks, training does not rely on epochs. Instead, the number of boosting rounds (n\_estimators) determines the size of the model. Cross-validation was not used because the data is temporal. A chronological split (80% train, 20% test) ensures that the model only learns from past matches and does not leak information from the future.

### 4.2 Performance Evaluation

Model + calibration	Accuracy	Log-loss	Mean-Brier
Logistic + Platt	0.61607	0.93570	0.18418
RandomForest + Platt	0.61607	0.91112	0.17709
XGBoost + raw	0.61607	0.90774	0.17550

Table 1: Model performance comparison

Outcome	RMSE (model vs book)	MAE (model vs book)
H	0.20713	0.16992
D	0.16699	0.13507
A	0.17963	0.14796

Table 2: Comparison between the best model (XGBoost raw) and bookmakers' implied probabilities

### 4.3 Visualizations

#### 4.3.1 Model vs bookmakers

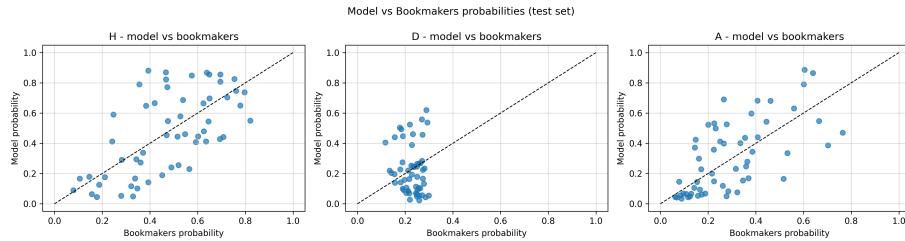


Figure 1: Scatter plots comparing the model's predicted probabilities with bookmaker implied probabilities for Home, Draw and Away outcomes.

The diagonal line represents perfect agreement. Most points lie reasonably close to this line, indicating that the model globally follows the structure of bookmaker pricing, while visible deviations highlight where the model systematically assigns higher or lower probabilities than the market.

#### 4.3.2 Error distribution

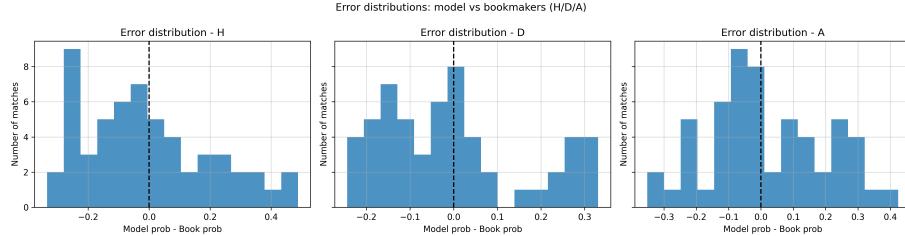
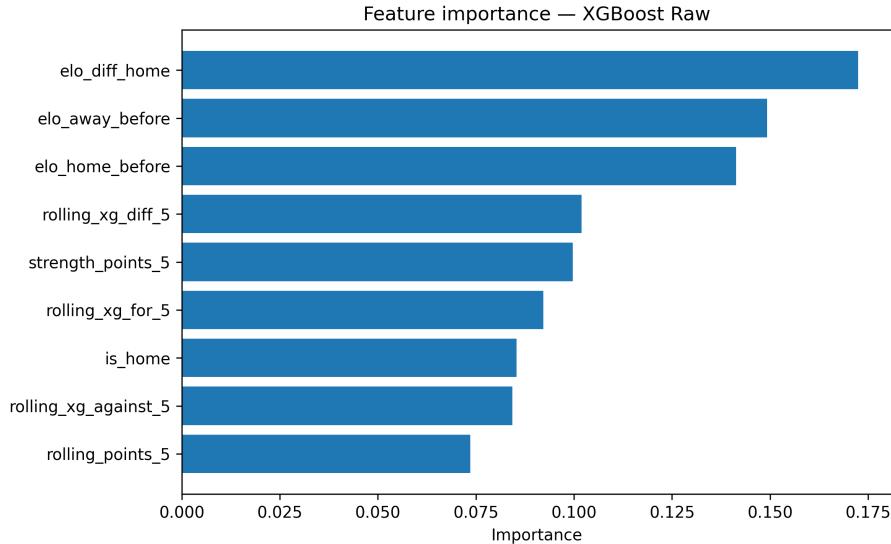


Figure 2: Distribution of prediction errors (model probability – bookmaker probability) for each outcome.

A distribution centered around zero indicates little systematic bias. Home and Draw errors are slightly shifted below zero, suggesting a small underestimation by the model, while Away errors are shifted above zero, indicating a small

overestimation. These distributions help quantify how the model diverges from bookmaker on average.

#### 4.3.3 Feature importance



*Figure 3: Feature importance for the XGBoost model (raw)*

This plot highlights which inputs the model relies on most when predicting match outcomes. ELO ratings dominate the decision process, indicating that team strength is the primary driver of predictions, while recent form (rolling xG and points) plays a secondary but meaningful role.

## 5. Discussion

Overall, the global pipeline performed well. The features engineered, particularly Elo ratings and the rolling xG metrics, significantly contributed to the results as the feature-importance plot demonstrated. The merging phase was quite demanding at first. As well as the rolling features construction where careful attention was required to avoid data leakage. Hyperparameter tuning for Random Forest and XGBoost was challenging but necessary to obtain well fitted model. Before running the models, it was expected that the overall performance of bookmakers would be better than that of the machine-learning models because of their richer information. These results confirm this partially: the probabilities of the bookmakers are still smoother and overall better-calibrated. However, the model performed better than expected regarding raw accuracy and could produce competitive probability estimates, especially concerning away wins. The relatively low difference in RMSE and MAE between XGBoost and

the bookmakers suggests that, even with only historical performance metrics, a data-driven approach can approximate bookmaker-level predictions. However, as the datasets are only based on one season this does not give enough information to the models to learn stable and long-term patterns. Moreover, the data is statistical, meaning the models do not have access to the pre-match crucial information (e.g. injuries, starting lineups etc) and context information (e.g. last game of the season is a must-win for certain teams). This missing information naturally creates a gap between the probabilities of the model and the one of the bookmakers. Surprisingly, the model systematically underestimates teams like Chelsea and Manchester City. For Chelsea, this is explained by the fact that they were inconstant during that season (finishing 12th), making it difficult for rolling features to understand the real level of the team. As for Manchester City, this is caused by the fact that Bookmakers constantly tend to “boost” them by assigning them high probabilities. They know how dominant the team is historically and how strong they become in the final stretch of the season, where the model only sees the data and does not know that they win the last 2 seasons.

## 6. Conclusion

### 6.1 Summary

The aim of this project was to investigate to what extent machine learning models can provide a probability as close as that of bookmakers’ implied probability, and how well football performance data alone can forecast match outcome probability? The outcome of this project is encouraging with 61% accuracy, average root means square error of 0.1846, and average mean absolute error of 0.1510 for the best model. We can conclude that football performance stats are helpful to get probabilities close to those of bookmakers. Yet the model, as expected, has revealed to be inaccurate for some games, this is caused by the lack of information. Finally, this study is an undeniable good start for comprehending the predictive potential and limits of data-driven modeling in football.

### 6.2 Future Work

Several additions could improve the results obtained. Most obvious one is to feed the models with more data. By adding more seasons of English Premier League, the models could learn more patterns and build a better historical understanding of each team. Adding informational features such as key players missing would likely have an impact as well. Extending the project to other championships could be interesting to see if some championships are “easier” to predict than others. Finally, implementing a betting back testing system would allow us to evaluate whether the models would have won or lost money over time. This would be for information purposes only because as experts often say: “The odds are never wrong; it’s the football that contradicts them”.

## References

1. Rathke, A. (2017). An examination of expected goals and shot efficiency in soccer. *Journal of Sports Analytics*. <https://www.redalyc.org/pdf/3010/301052437005.pdf>
2. Mead, A., Pollard, R., & Reep, C. (2023). Expected goals in football: Improving model performance and demonstrating value. *PLOS ONE*. <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0282295>
3. Hvattum, L. M., & Arntzen, H. (2010). Using ELO ratings for match result prediction in association football. *International Journal of Forecasting*, 26(3), 460–470. <https://www.sciencedirect.com/science/article/pii/S0169207009001708?via%3Dhub>
4. Tax, N., & Joustra, Y. (2015). Predicting the Dutch football competition using public data: A machine learning approach. [https://www.researchgate.net/publication/282026611\\_Predicting\\_The\\_Dutch\\_Football\\_Competition](https://www.researchgate.net/publication/282026611_Predicting_The_Dutch_Football_Competition)
5. Niculescu-Mizil, A., & Caruana, R. (2005). Predicting Good Probabilities with Supervised Learning. <https://www.cs.cornell.edu/~alexn/papers/calibration.icml05.crc.rev3.pdf>
6. Levitt, S. D. (2004). Why are gambling markets organised so differently from financial markets? *The Economic Journal*, 114. <https://pricetheory.uchicago.edu/levitt/Papers/LevittWhyAreGamblingMarkets2004.pdf>
7. Clarke, S. R. (2017). Adjusting Bookmaker's Odds to Allow for Overround [https://www.researchgate.net/publication/326510904\\_Adjusting\\_Bookmaker%27s\\_Odds\\_to\\_Allow\\_for\\_Overround](https://www.researchgate.net/publication/326510904_Adjusting_Bookmaker%27s_Odds_to_Allow_for_Overround)
8. Football-Data.co.uk (2023). English Premier League betting odds dataset (2022–2023 season). Available at: <https://www.football-data.co.uk/englandm.php> (accessed 2025).
9. Dataquestio (2023). FBref football datasets. Available at: [https://github.com/dataquestio/project-walkthroughs/blob/master/football\\_matches/matches.csv](https://github.com/dataquestio/project-walkthroughs/blob/master/football_matches/matches.csv) (accessed 2025).

## Appendices

### Appendix A: Additional Results





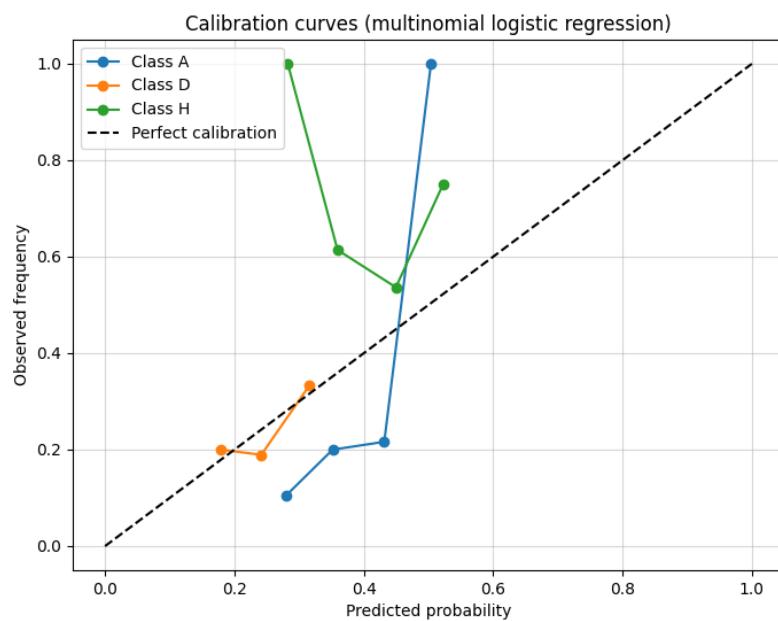
Figure 4: Reliability curves comparing the model's predicted probabilities with bookmakers' implied probabilities for Home, Draw, and Away outcomes.

This figure illustrates how closely the predicted probabilities of the model and bookmakers align with the observed outcome frequencies, highlighting differences in calibration across outcomes and probability ranges.




*Figure 5: Calibration plot of the best performing model (XGBoost raw)*

This calibration curve illustrates how the predicted probabilities of the raw XG-Boost model (best model based on log-loss) compare to the observed outcome frequencies, highlighting deviations from perfect calibration (diagonal line) across home, draw, and away outcomes

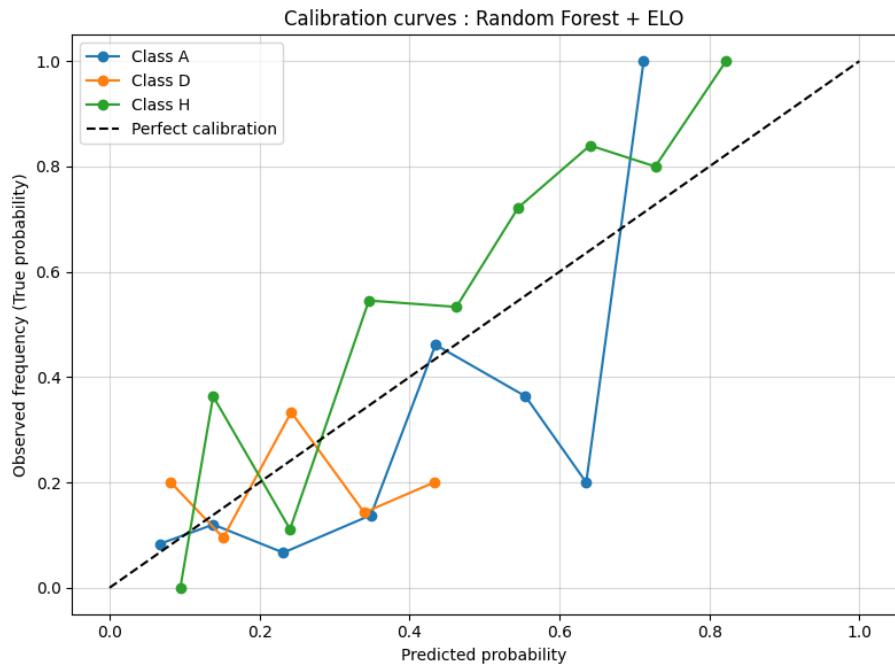


*Figure 6: Calibration plot of the baseline model without elo-based features*



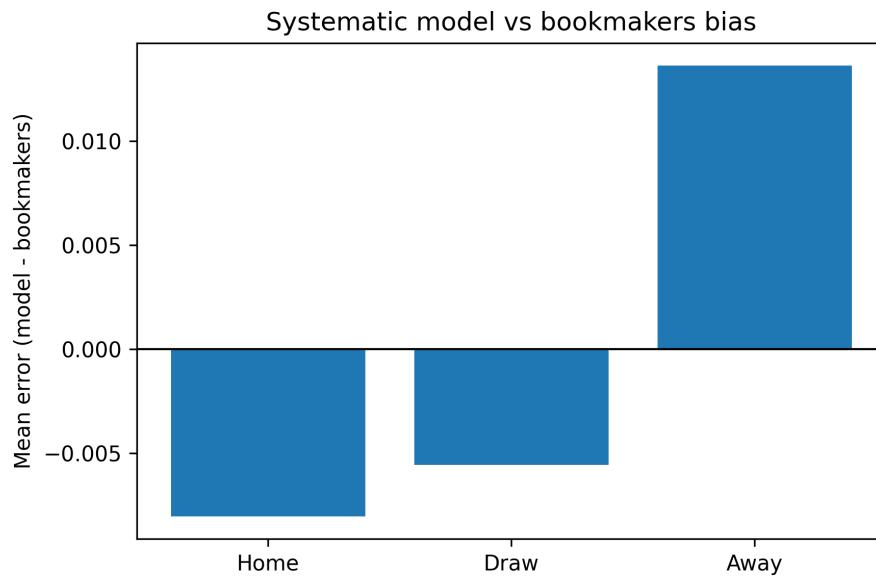
*Figure 7: Calibration plot of the baseline model with elo-based features (raw)*

This plot illustrates the impact of introducing Elo-based features, which already improves calibration compared to the baseline model



*Figure 8: Calibration plot of the Random forest model (raw)*

The Random Forest calibration curve shows a better alignment with observed frequencies than the baseline models, although some miscalibration remains, especially for extreme probability values.



*Figure 9: Average difference between model-implied probabilities and bookmaker-implied probabilities for each outcome.*

Negative values indicate that the model underestimates probabilities relative to bookmakers, while positive values indicate overestimation, revealing systematic biases across Home, Draw, and Away outcomes.

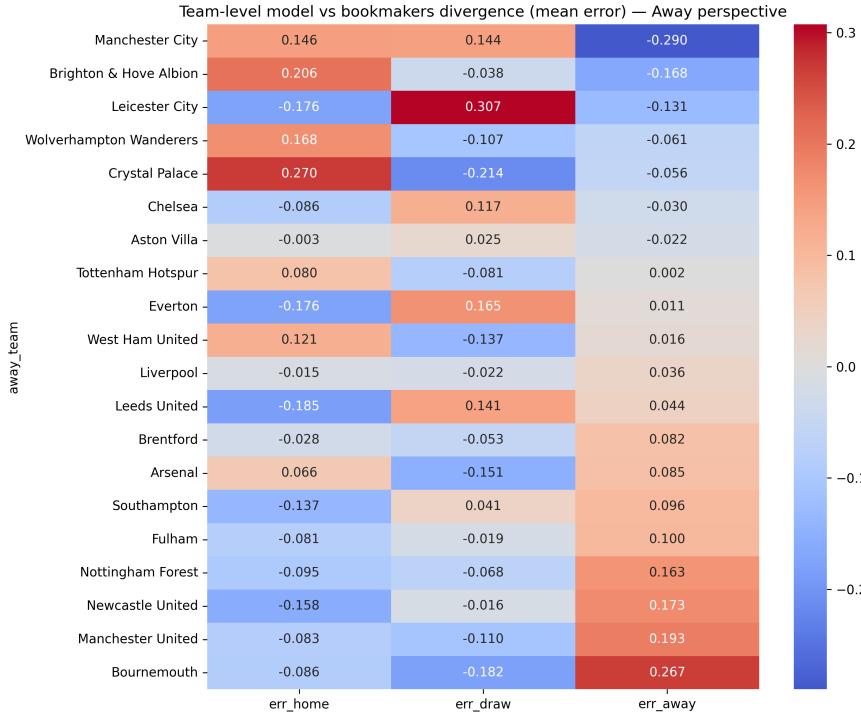


Figure 10: Team-level model vs bookmaker divergence (mean error) for away matches.

This heatmap displays the average difference between the model’s predicted probabilities and bookmakers’ implied probabilities from an away-team perspective (model – bookmaker). Positive values indicate that the model assigns a higher probability to away wins than bookmakers, while negative values indicate underestimation.

A notable example is Manchester City, for which the model strongly underestimates away win probabilities compared to bookmakers. This suggests that the model fails to fully capture Manchester City’s exceptional dominance, particularly in away fixtures, relying only on historical performance metrics. In contrast, bookmakers appear to incorporate additional contextual information and long-term dominance, leading to systematically higher away-win probabilities for such elite teams. This highlights a structural limitation of purely data-driven models when faced with historically dominant clubs.

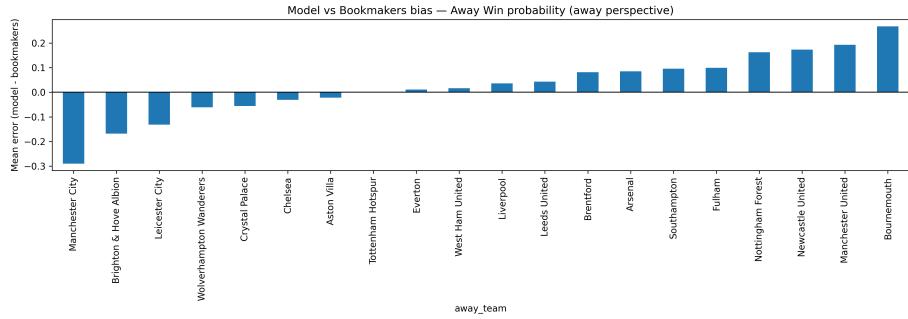


Figure 11: Bar plot representation of team-level away bias.

This bar plot provides a more intuitive visualization of the same away-team biases shown in Figure 10 and is included for completeness.

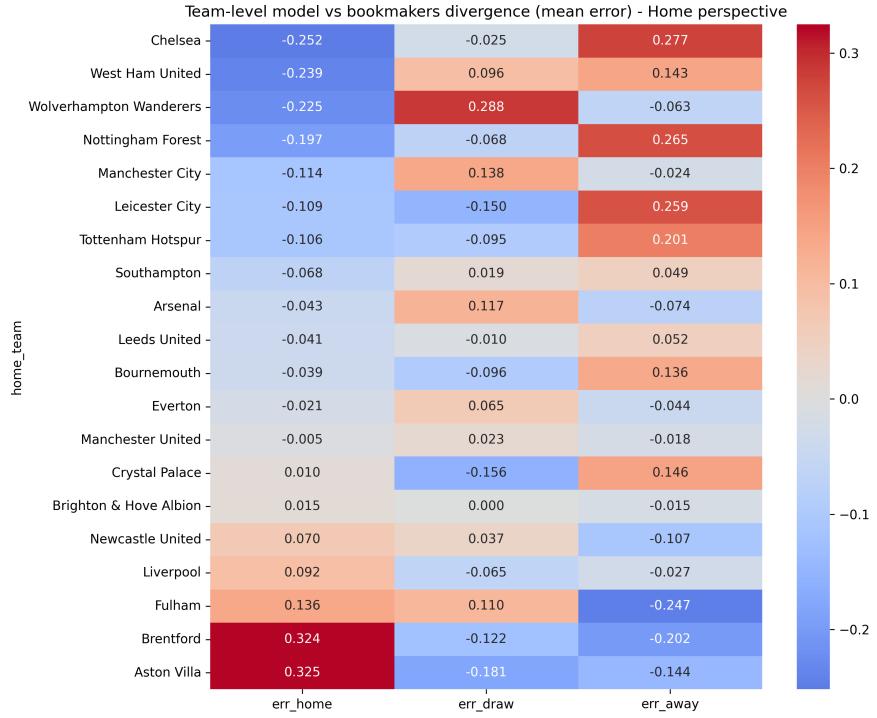


Figure 12: Team-level model vs bookmaker divergence (mean error) for home matches.




Figure 13: Bar plot representation of team-level home bias.




Figure 14: Distribution of probabilities best model vs bookmakers

It shows the distribution of predicted probabilities for Home, Draw and Away outcomes.

## Appendix B: Code Repository

**GitHub Repository:** <https://github.com/MattisD3/Data-Science-in-Football-Comparing-Model-Predictions-with-Betting-Market-Odds>

## Appendix C: Ai tools used

ChatGPT (OpenAI) was used as a learning and support tool during this project. It mainly helped to: - better understand machine learning concepts that were new to me, such as model calibration and hyperparameter tuning, - clarify the behavior of different models (Logistic Regression, Random Forest, XGBoost), - provide guidance when debugging existing code.

All data processing, modeling decisions, experiments, and interpretations were carried out by me. ChatGPT was used in a similar way to online documentation or forums, as an explanatory and educational resource.

### Repository Structure

```
DATA-SCIENCE-IN-FOOTBALL-COMPARING-MODEL-PREDICTIONS-WITH-BETTING-MARKET-ODDS/
  README.md
  requirements.txt
  environment.yml
  project_report.md
  PROPOSAL
  main.py
  data/
    raw/
      matches.csv
      data_bookmakers_22_23.csv
    processed/
      clean_bookmakers_22_23.csv
      clean_matches_22_23.csv
      elo_ratings_22_23.csv
```

```

        features_matches_long_22_23.csv
        features_matches_long Elo_22_23.csv
        matches_long_22_23.csv
        matches_wides_22_23.csv

notebooks/

src/
    A_clean_matches.py
    B_clean_bookmakers_data.py
    C_merge_data.py
    D_feature_engineering.py
    E_model_baseline_logistic.py
    F_elastic_net.py
    G_add_elastic_net_to_long_features.py
    H_model_baseline_logistic_with_elastic_net.py
    I_random_forest.py
    J_xgboost.py
    K_calibration_experiments.py
    L_model_vs_bookmakers.py

results/
    calibration_model_vs_book_22_23.png
    calibration_plot_baseline.png
    calibration_plot_best.png
    calibration_plot_elastic_net.png
    calibration_plot_random_forest.png
    calibration_plot_xgboost.png
    calibration_summary_22_23.csv
    errors_model_vs_book_22_23.png
    feature_importance_xgb_raw_22_23.png
    global_bias_bar_22_23.png
    model_vs_bookmakers_22_23.csv
    prob_dist_model_vs_book_22_23.png
    scatter_model_vs_book_22_23.png
    team_away_bias_bar_22_23.png
    team_away_bias_heatmap_22_23.png
    team_home_bias_bar_22_23.png
    team_home_bias_heatmap_22_23.png

|
tests/
docs/
.gitignore

```

### **Installation Instructions**

```
git clone https://github.com/MattisD3/Data-Science-in-Football  
-Comparing-Model-Predictions-with-Betting-Market-Odds  
cd Data-Science-in-Football-Comparing-Model-Predictions-with-Betting-Market-Odds  
conda env create -f environment.yml  
conda activate football-outcome-prediction
```

### **Reproducing Results**

```
python main.py
```